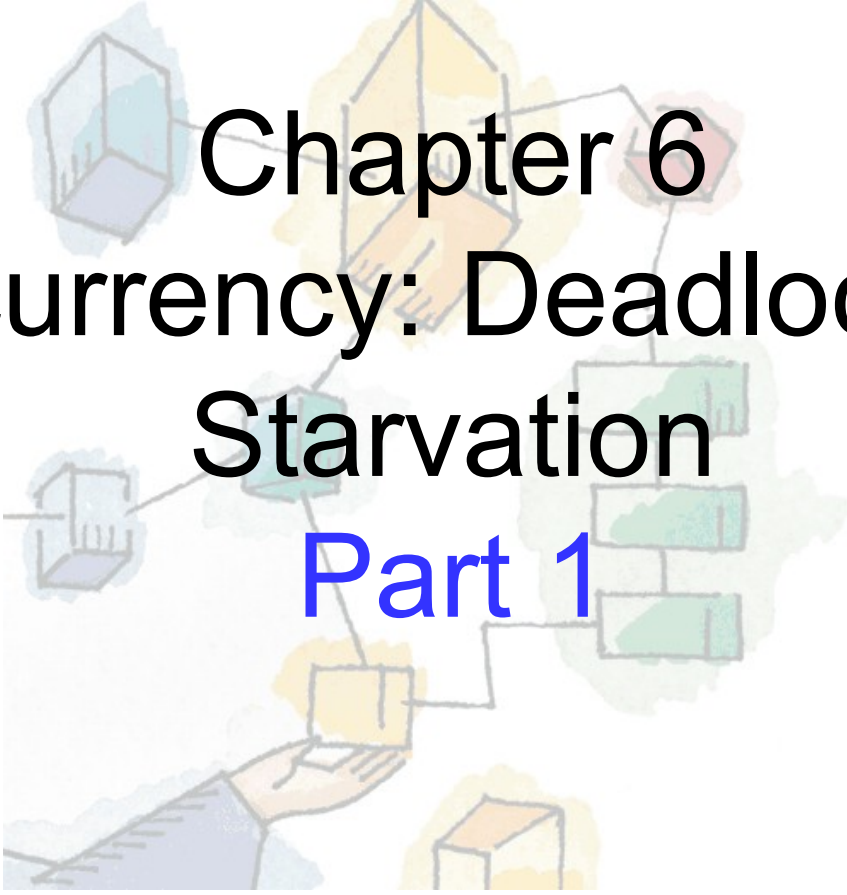


*Operating Systems:  
Internals and Design Principles, 8/E*  
William Stallings



# Chapter 6

## Concurrency: Deadlock and Starvation

### Part 1

Patricia Roy  
Manatee Community College, Venice, FL  
©2015, Prentice Hall

# Outline


- Principle of deadlock
- Deadlock Prevention
- Deadlock avoidance
- Deadlock Detection
- Integrated Deadlock Strategy
- Dining philosophers Problem
- UNIX concurrency Mechanisms
- Linux Kernel Concurrency Mechanism

# The Anology

*When two trains approach each other at a crossing, both shall come to a full stop and neither shall start up again until the other has gone. Statute passed by the Kansas State Legislature, early in the 20th century.*

—*A TREASURY OF RAILROAD FOLKLORE*,  
B. A. Botkin and Alvin F. Harlow

# Deadlock



- **Permanent blocking of a set of processes** that either compete for system resources or communicate with each other
- A set of processes is deadlocked when each process in the set is **blocked awaiting an event** that can only be triggered by another blocked process in the set
- Unlike other problem, there is **No efficient solution**
- Involve conflicting needs for resources by two or more processes







# Deadlock

Resources are available

Ignore the rules and proceed



(a) Deadlock possible



(b) Deadlock

All resources are seized

Figure 6.1 Illustration of Deadlock

# Deadlock

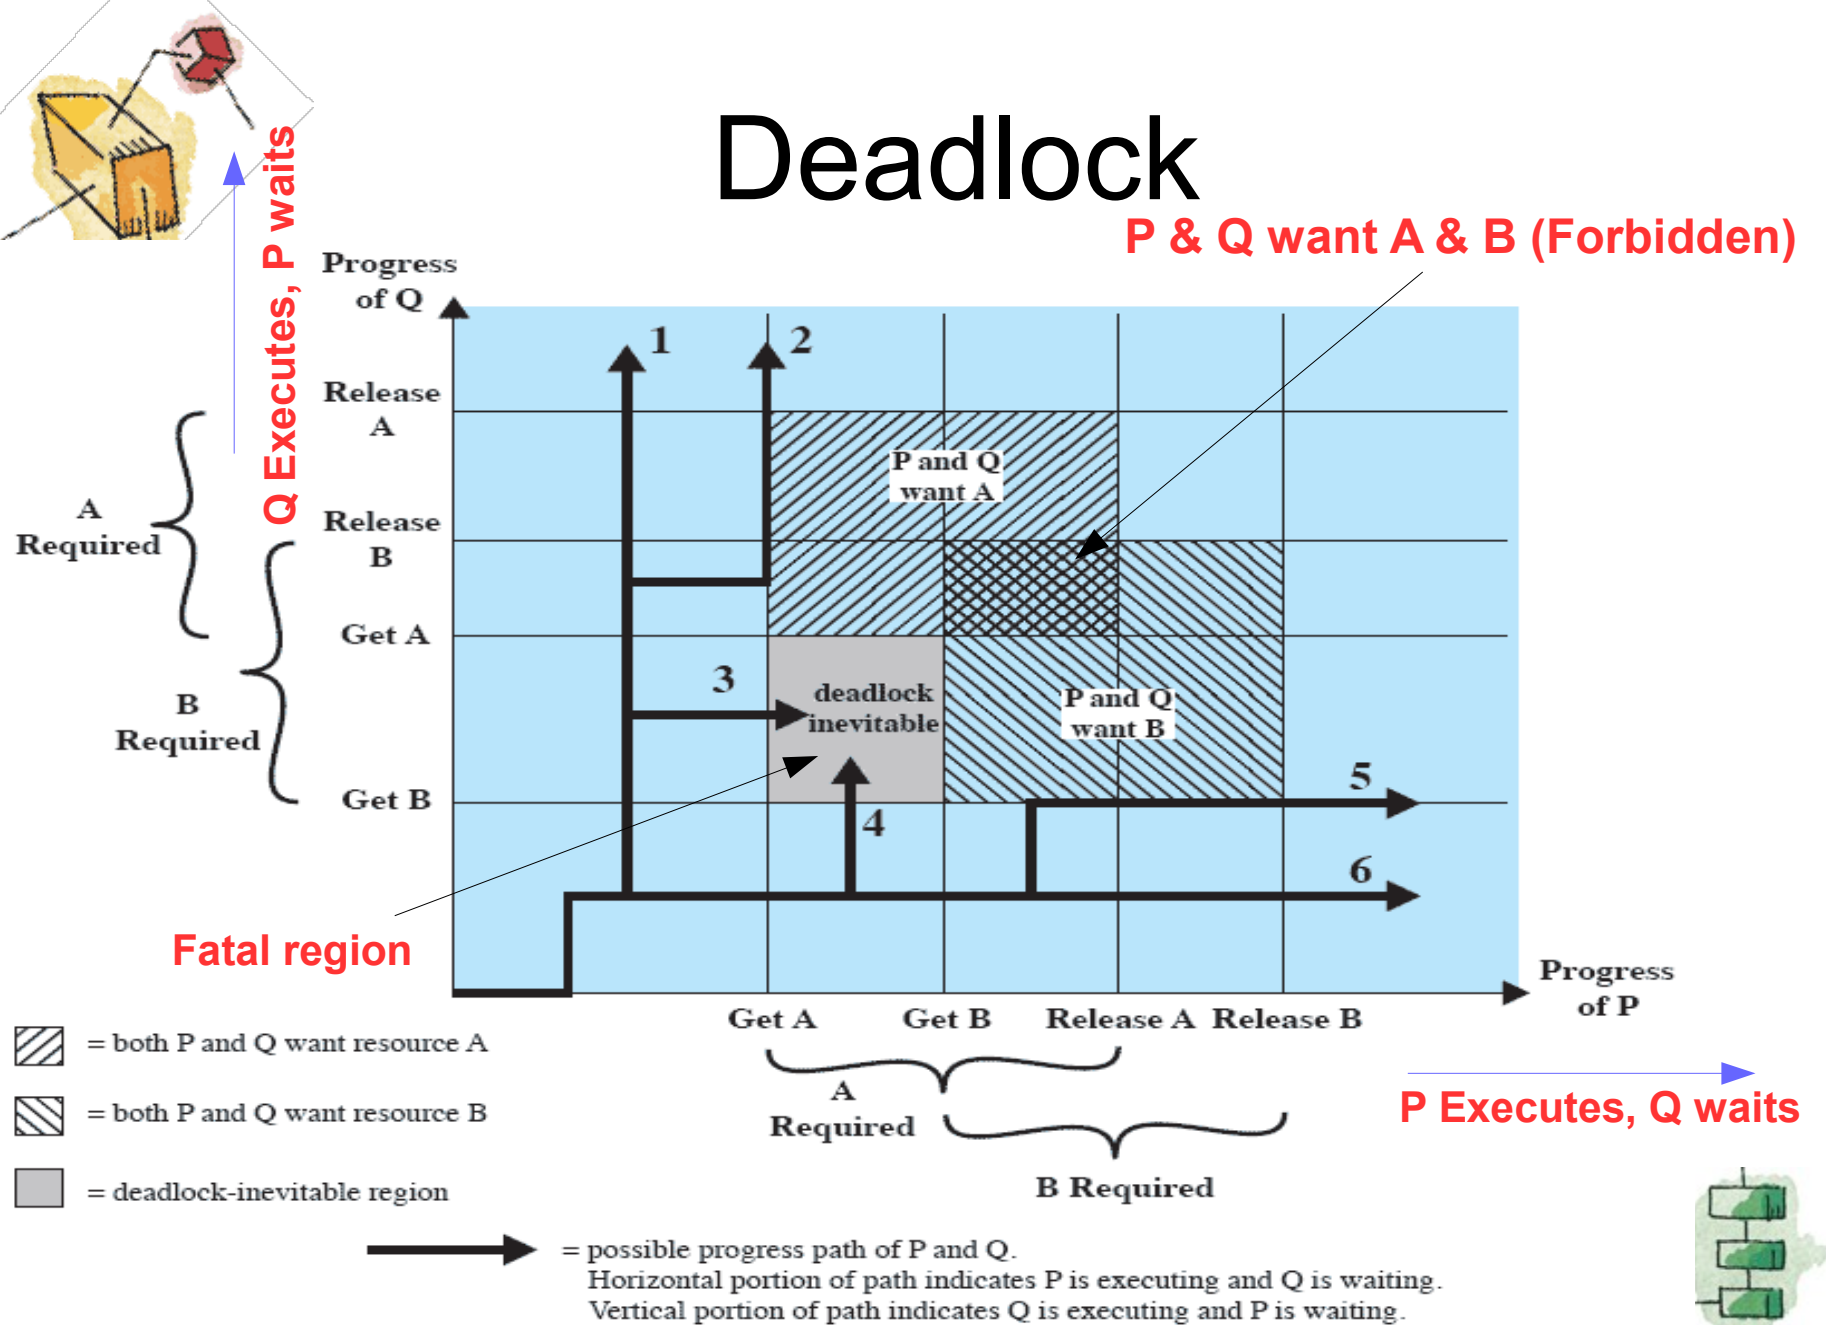


Figure 6.2 Example of Deadlock



# 6 different execution paths

- **Q acquires B and then A and then releases B and A.** When P resumes execution, it will be able to acquire both resources
- **Q acquires B and then A.** P executes and blocks on a request for A. Q releases B and A. When P resumes execution, it will be able to acquire both resources
- **Q acquires B and then P acquires A.** Deadlock is inevitable because as execution proceeds, Q will block on A and P will block on B
- **P acquires A and then Q acquires B.** Deadlock is inevitable because as execution proceeds, Q will block on A and P will block on B
- **P acquires A and then B.** Q executes and blocks on a request for B. P releases A and B. When Q resumes execution, it will be able to acquire both resources
- **P acquires A and then B and then releases A and B** When Q resumes execution, it will be able to acquire both resources



# Deadlock





Figure 6.3 Example of No Deadlock [BACO03]






# Reusable Resources

- Used by only one process at a time and not depleted by that use
- Processes obtain resources that they later release for reuse by other processes
- E.g: Processors, I/O channels, main and secondary memory, devices, and data structures such as files, databases, and semaphores
- Deadlock occurs if each process holds one resource and requests the other





# Reusable Resources

## Process P

Step	Action
$p_0$	Request (D)
$p_1$	Lock (D)
$p_2$	Request (T)
$p_3$	Lock (T)
$p_4$	Perform function
$p_5$	Unlock (D)
$p_6$	Unlock (T)

## Process Q


Step	Action
$q_0$	Request (T)
$q_1$	Lock (T)
$q_2$	Request (D)
$q_3$	Lock (D)
$q_4$	Perform function
$q_5$	Unlock (T)
$q_6$	Unlock (D)

**D = Disk File**

**T = Tape Drive**



Figure 6.4 Example of Two Processes Competing for Reusable Resources






# Reusable Resources

- Space is available for allocation of 200Kbytes, and the following sequence of events occur



- Deadlock occurs if both processes progress to their second request






# Consumable Resources



- Created (produced) and destroyed (consumed)
- When acquired by a process, it ceases to exist
- E.g: Interrupts, signals, messages, and information in I/O buffers






# Example of Deadlock

- Each process attempts to receive a message from the other process, then send a message to the other one.




- Deadlock may occur if a Receive message is blocking (until the message is received)
- Design error – cause deadlock, difficult to detect
- May take a rare combination of events to cause deadlock





Approach	Resource Allocation Policy	Different Schemes	Major Advantages	Major Disadvantages
Prevention	Conservative; undercommits resources	Requesting all resources at once	<ul style="list-style-type: none"> <li>• Works well for processes that perform a single burst of activity</li> <li>• No preemption necessary</li> </ul>	<ul style="list-style-type: none"> <li>• Inefficient</li> <li>• Delays process initiation</li> <li>• Future resource requirements must be known by processes</li> </ul>
		Preemption	<ul style="list-style-type: none"> <li>• Convenient when applied to resources whose state can be saved and restored easily</li> </ul>	<ul style="list-style-type: none"> <li>• Preempts more often than necessary</li> </ul>
		Resource ordering	<ul style="list-style-type: none"> <li>• Feasible to enforce via compile-time checks</li> <li>• Needs no run-time computation since problem is solved in system design</li> </ul>	<ul style="list-style-type: none"> <li>• Disallows incremental resource requests</li> </ul>
Avoidance	Midway between that of detection and prevention	Manipulate to find at least one safe path	<ul style="list-style-type: none"> <li>• No preemption necessary</li> </ul>	<ul style="list-style-type: none"> <li>• Future resource requirements must be known by OS</li> <li>• Processes can be blocked for long periods</li> </ul>
Detection	Very liberal; requested resources are granted where possible	Invoke periodically to test for deadlock	<ul style="list-style-type: none"> <li>• Never delays process initiation</li> <li>• Facilitates online handling</li> </ul>	<ul style="list-style-type: none"> <li>• Inherent preemption losses</li> </ul>





# Resource Allocation Graphs


- Directed graph that depicts a state of the system of resources and processes



(a) Resource is requested




(b) Resource is held




(c) Circular wait


Deadlock




(d) No deadlock

Many resources






# Resource Allocation Graphs



**Figure 6.6** Resource Allocation Graph for Figure 6.1b








# Conditions for Deadlock

Necessary, but not sufficient


- **Mutual exclusion**
  - Only one process may use a resource at a time
- **Hold-and-wait**
  - A process may hold allocated resources while awaiting assignment of others
- **No preemption**
  - No resource can be forcibly removed from a process holding it
- **Circular wait**
  - A closed chain of processes exists, such that each process holds at least one resource needed by the next process in the chain



For a deadlock to actually take place




Possibility of Deadlock	Existence of Deadlock
Mutual Exclusion	Mutual Exclusion
No preemption	No preemption
Hold and wait	Hold and wait
	<b>Circular Wait</b>



Unresolvable




# Deadlock Prevention



- **Mutual Exclusion (ME)**
  - If access to resource require ME, it must be supported by the OS.  
Eg: file: multiple read, single write
- **Hold and Wait**
  - Require a process request all of its required resources at one time.  
**Inefficient**: longtime waiting for all resources, remain unused, denied to other processes.
- **No Preemption**
  - First: if a process holding certain resources is denied a further request, that process must release its original resources and request again
  - If the resources are held by others, OS may preempt the 2nd process to require it releases its resources (**if no 2 processes possessed the same priority**)
- **Circular Wait**
  - Define a linear ordering of resource types






# Deadlock Avoidance

- Differs subtly from deadlock prevention
- Allows the 3 necessary conditions, but make intelligent choices to assure that deadlock point is never reached
  - *Is the current resource alloc. will lead to potential deadlock?*
- Thus, it allows more concurrency than prevention
- A decision is made dynamically whether the current resource allocation request will, if granted, potentially lead to a deadlock
- Thus, requires knowledge of future process requests





# Two Approaches to Deadlock Avoidance

- Do not start a process if its demands might lead to deadlock
- Do not grant an incremental resource request to a process if this allocation might lead to deadlock

