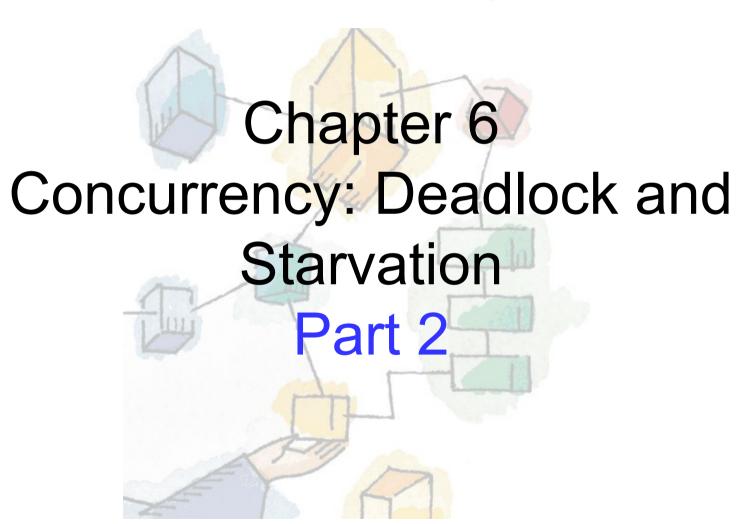
Operating Systems: Internals and Design Principles, 8/E William Stallings



Patricia Roy
Manatee Community College, Venice, FL
©2015, Prentice Hall

Resource Allocation Denial

- The strategy referred to as the banker's algorithm
- State of the system is the current allocation of resources to processes:
 - Vectors: Resource & Available
 - Matrices: Claim & Allocation
- Safe state is where there is at least one sequence of resource allocation that does not result in deadlock
- Unsafe state is a state that is not safe





Deadlock Avoidance

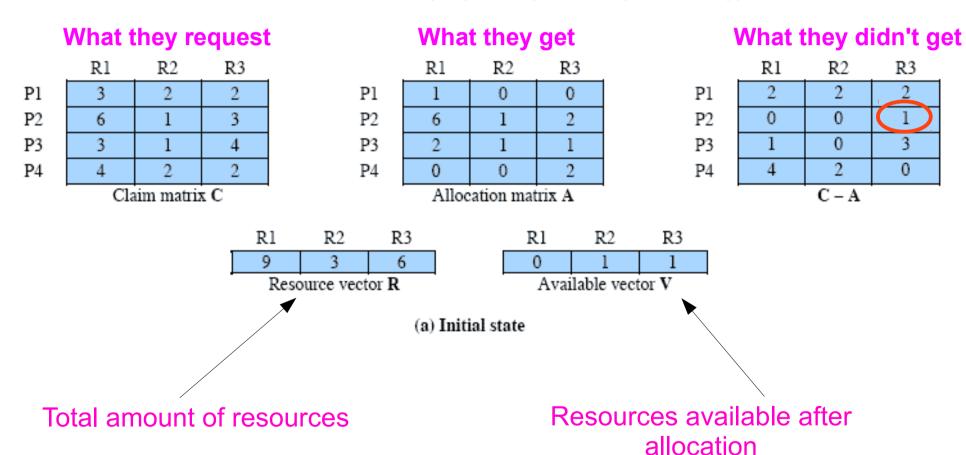
- Differs subtly from deadlock prevention
- Allows the 3 necessary conditions, but make intelligent choices to assure that deadlock point is never reached
 - Is the current resource alloc. will lead to potential deadlock?
- Thus, it allows more concurrency than prevention
- A decision is made dynamically whether the current resource allocation request will, if granted, potentially lead to a deadlock
- Thus, requires knowledge of future process requests





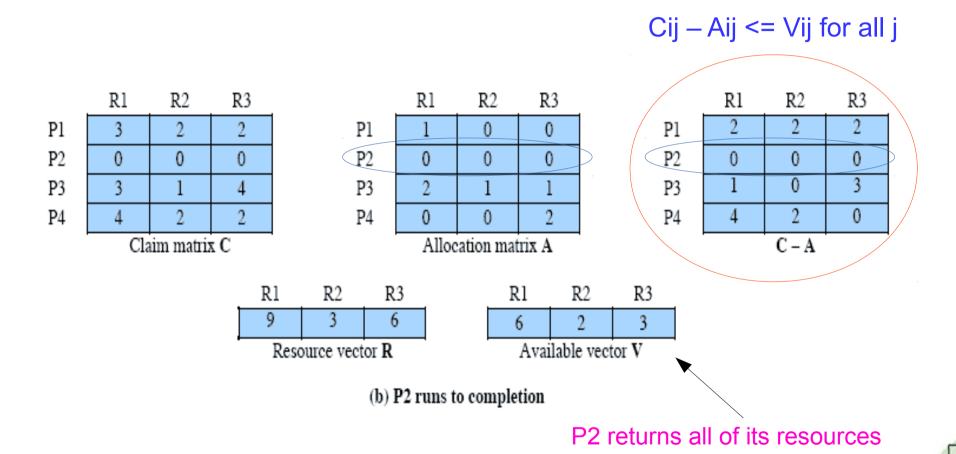
Determination of a Safe State

Can any of the 4 processes be run to completion with the resources available? (Cij – Aij <= Vij for all j)





Is this a safe state?



** Now, each of the remaining processes could be completed



Is this a safe state?

Suppose now P1 runs to completion, and returns all of its resources

| | Rl | R2 | R3 | |
|----|----------------|----|----|--|
| Pl | 0 | 0 | 0 | |
| P2 | 0 | 0 | 0 | |
| P3 | 3 | l | 4 | |
| P4 | 4 | 2 | 2 | |
| | Claim matrix C | | | |

| | Rl | R2 | R3 | |
|----|---------------------|----|----|--|
| Pl | 0 | 0 | 0 | |
| P2 | 0 | 0 | 0 | |
| P3 | 2 | l | l | |
| P4 | 0 | 0 | 2 | |
| , | Allocation matrix A | | | |

| | R1 | R2 | R3 |
|----|----|-------|----|
| Pl | 0 | 0 | 0 |
| P2 | 0 | 0 | 0 |
| P3 | l | 0 | 3 |
| P4 | 4 | 2 | 0 |
| | | C – A | |

| Rl | R2 | R3 | |
|-------------------|----|----|--|
| 9 | 3 | 6 | |
| Resource vector R | | | |

| R1 | R2 | R3 | |
|--------------------|----|----|--|
| 7 | 2 | 3 | |
| Available vector V | | | |

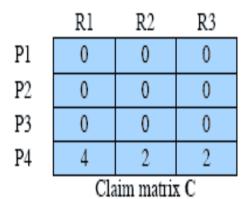
(c) P1 runs to completion





Is this a safe state?

Suppose now P3 runs to completion, and returns all of its resources



| | Rl | R2 | R3 |
|---------------------|----|----|----|
| Pl | 0 | 0 | 0 |
| P2 | 0 | 0 | 0 |
| P3 | 0 | 0 | 0 |
| P4 | 0 | 0 | 2 |
| Allocation matrix A | | | |

| | Rl | R2 | R3 |
|----|----|-------|----|
| Pl | 0 | 0 | 0 |
| P2 | 0 | 0 | 0 |
| P3 | 0 | 0 | 0 |
| P4 | 4 | 2 | 0 |
| | | C – A | |

| | Rl | R2 | R3 | |
|---|-------------------|----|----|--|
| | 9 | 3 | 6 | |
| • | Resource vector R | | | |

| | Rl | R2 | R3 | |
|--------------------|----|----|----|--|
| | 9 | 3 | 4 | |
| Available vector V | | | | |

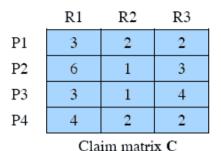
(d) P3 runs to completion

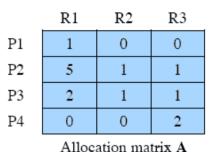
*** Therefore, all of these processes are in the safe state





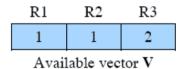
Determination of an Unsafe State



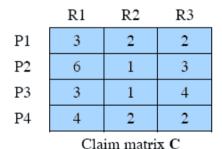


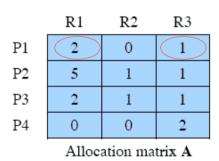
| | R1 | R2 | R3 |
|----|----|-------|----|
| P1 | 2 | 2 | 2 |
| P2 | 1 | 0 | 2 |
| P3 | 1 | 0 | 3 |
| P4 | 4 | 2 | 0 |
| | | C – A | |

| | R1 | R2 | R3 |
|---|-------------------|----|----|
| | 9 | 3 | 6 |
| _ | Resource vector R | | |



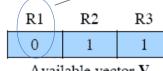
(a) Initial state





| | R1 | R2 | R3 |
|----|-----|-------|----|
| P1 | 1 | 2 | 1 |
| P2 | 1 | 0 | 2 |
| P3 | 1 | 0 | 3 |
| P4 | 4 / | 2 | 0 |
| | | C – A | |

| . R1 | | R2 | R3 |
|-------------------|--|----|----|
| 9 | | 3 | 6 |
| Resource vector R | | | |



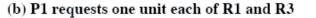
0 1 1 not a safe state!

Available vector V Potential for DEADLOCK.

P1 request should be denied,

P1 should be blocked!

Not possible, thus this is





Deadlock Avoidance Logic

```
struct state {
    int resource[m];
    int available[m];
    int claim[n][m];
    int alloc[n][m];
}
```

(a) global data structures





Deadlock Detection

- Deadlock prevention strategies are very conservative
 - Solve deadlock by limiting access to resources
 - Imposing restrictions on processes

- Deadlock detection do not limit resource access
 - Requested resources are granted to processes whenever possible
 - The OS will periodically performs an algorithm to detect the circular wait condition





Deadlock Detection

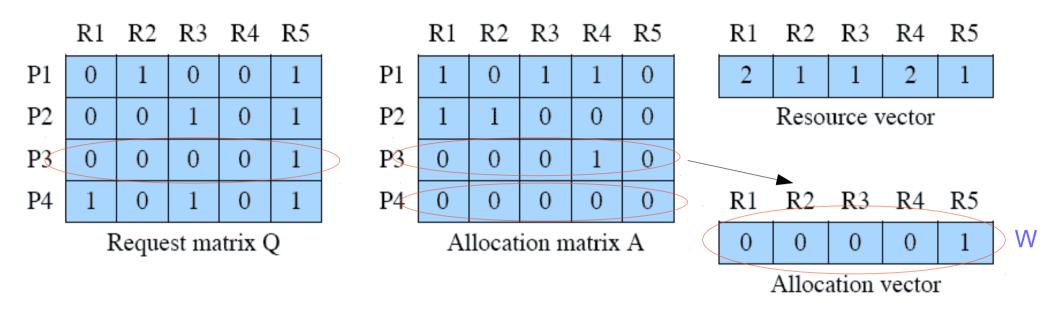


Figure 6.10 Example for Deadlock Detection

$$W = W + (00010) = (00011)$$

- 1. Mark each process that has a row in Allocation matrix of all zeros
- 2. Initialise a temporary vector W to equal the Available vector
- 3. Find an index *i* such that process *i* is currently unmarked and the *ith* row of Q is less than or equal to W. (Q*ik* <= W*k* for 1<=k<=m). If no such row is found, terminate the algorithm (DEADLOCK)
- 4. If such a row is found, mark process i and add the corresponding row of the allocation matrix to W. (Wk = Wk + Aik for 1<=k<=m). Return to step 3.

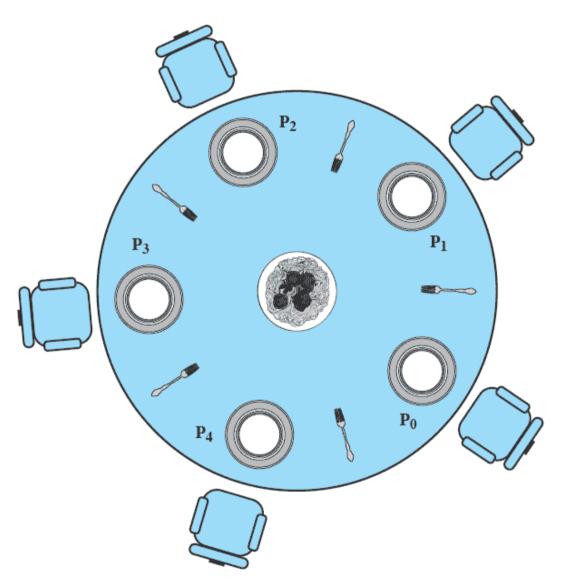




Strategies Once Deadlock Detected

- Abort all deadlocked processes
- Back up each deadlocked process to some previously defined checkpoint, and restart all process
 - Risk: Original deadlock may recur
- Successively abort deadlocked processes until deadlock no longer exists.
 - order depends on the minimum cost
- Successively preempt resources until deadlock no longer exists.
 - Also use minimum cost-based approach

Dining Philosophers Problem



5 philosophers in a house

Thinking & Eating!

The only food that contribute: Spaghetti!

2 forks each to eat spaghetti

Rule:

- 1. Satisfy mutual exclusion
 - No 2 philosophers can use the same forks at the same time
- 2. Avoid deadlock and starvation



Figure 6.11 Dining Arrangement for Philosophers

Solutions...

```
/* program
               diningphilosophers */
semaphore fork [5] = {1};
int i;
void philosopher (int i)
     while (true) {
          think();
          wait (fork[i]);
          wait (fork [(i+1) mod 5]);
          eat();
          signal(fork [(i+1) mod 5]);
          signal(fork[i]);
void main()
     parbegin (philosopher (0), philosopher (1), philosopher
(2)_{r}
          philosopher (3), philosopher (4));
```

Figure 6.12 A First Solution to the Dining Philosophers Problem





Solutions...

```
/* program diningphilosophers */
semaphore fork[5] = {1};
semaphore room = {4};
int i;
void philosopher (int i)
   while (true) {
    think();
    wait (room);
    wait (fork[i]);
     wait (fork [(i+1) mod 5]);
     eat();
     signal (fork [(i+1) mod 5]);
     signal (fork[i]);
     signal (room);
void main()
   parbegin (philosopher (0), philosopher (1), philosopher (2),
          philosopher (3), philosopher (4));
```

Figure 6.13 A Second Solution to the Dining Philosophers Problem





UNIX Concurrency Mechanisms

Pipes

 Circular buffer allowing two processes to communicate on the producer-consumer model

Messages

- For processes to engage in msg passing mailbox
- Read from empty mailbox -> block
- Shared memory
 - Virtual memory shared by multiple processes
- Semaphores
 - SemWait & SemSignal
- Signals
 - Similar to hardware interrupt, but has no priorities





UNIX Signals

| Value | Name | Description |
|-------|---------|--|
| 01 | SIGHUP | Hang up; sent to process when kernel assumes that the user of that process is doing no useful work |
| 02 | SIGINT | Interrupt |
| 03 | SIGQUIT | Quit; sent by user to induce halting of process and production of core dump |
| 04 | SIGILL | Illegal instruction |
| 05 | SIGTRAP | Trace trap; triggers the execution of code for process tracing |
| 06 | SIGIOT | IOT instruction |
| 07 | SIGEMT | EMT instruction |
| 08 | SIGFPE | Floating-point exception |
| 09 | SIGKILL | Kill; terminate process |
| 10 | SIGBUS | Bus error |
| 11 | SIGSEGV | Segmentation violation; process attempts to access location outside its virtual address space |
| 12 | SIGSYS | Bad argument to system call |
| 13 | SIGPIPE | Write on a pipe that has no readers attached to it |
| 14 | SIGALRM | Alarm clock; issued when a process wishes to receive a signal after a period of time |
| 15 | SIGTERM | Software termination |
| 16 | SIGUSR1 | User-defined signal 1 |
| 17 | SIGUSR2 | User-defined signal 2 |
| 18 | SIGCHLD | Death of a child |
| 19 | SIGPWR | Power failure |





Linux Kernel Concurrency Mechanism

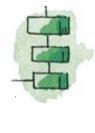
- Includes all the concurrency mechanisms found in UNIX (e.g. SVR4 including pipes, messages, shared memory, and signals)
- Three types:
 - Atomic operations
 - Spinlocks
 - Semaphores





Atomic operations

- Execute without interruption and without interference
- Use to avoid simple race conditions
- Simplest of the approaches to kernel synchronization
- Two Types:
 - Atomic integer operation
 - Atomic bitmap operation





Linux Atomic Operations

Table 6.3 Linux Atomic Operations

| Atomic Integer Operations | | | |
|--|---|--|--|
| ATOMIC_INIT (int i) | At declaration: initialize an atomic_t to i | | |
| int atomic_read(atomic_t *v) | Read integer value of v | | |
| <pre>void atomic_set(atomic_t *v, int i)</pre> | Set the value of v to integer i | | |
| void atomic_add(int i, atomic_t *v) | Additov | | |
| void atomic_sub(int i, atomic_t *v) | Subtract i from v | | |
| void atomic_inc(atomic_t *v) | Add 1 to v | | |
| void atomic_dec(atomic_t *v) | Subtract 1 from v | | |
| <pre>int atomic_sub_and_test(int i, atomic_t *v)</pre> | Subtract i from v; return 1 if the result is zero; return 0 otherwise | | |
| <pre>int atomic_add_negative(int i, atomic_t *v)</pre> | Add i to v; return 1 if the result is negative; return 0 otherwise (used for implementing semaphores) | | |
| int atomic_dec_and_test(atomic_t *v) | Subtract 1 from v; return 1 if the result is zero; return 0 otherwise | | |
| int atomic_inc_and_test(atomic_t *v) | Add 1 to v; return 1 if the result is zero; return 0 otherwise | | |





Spinlocks

- Most common technique for protecting a critical section in Linux
- Can only be acquired by one thread at a time
- Any other thread will keep trying (spinning) until it can acquire the lock
- Built on an integer location in memory that is checked by each thread before it enters its critical section
 - If the value is 0, the thread enters the critical region and set the value to
 - If the value is nonzero, the thread will continually check the value until the value is 0.
- Effective in situations where the wait time for acquiring a lock is expected to be very short
- Disadvantage:
 - locked-out threads continue to execute in a busy-waiting mode





Spinlocks

```
spin_lock (&lock)
  /*critical section */
spin_unlock (&lock)
```





Linux Spinlocks

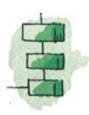
| void spin_lock(spinlock_t *lock) | Acquires the specified lock, spinning if needed until it is available |
|--|---|
| void spin_lock_irq(spinlock_t *lock) | Like spin_lock, but also disables interrupts on the local processor |
| <pre>void spin_lock_irqsave(spinlock_t *lock, unsigned long flags)</pre> | Like spin_lock_irq, but also saves the current interrupt state in flags |
| void spin_lock_bh(spinlock_t *lock) | Like spin_lock, but also disables the execution of all bottom halves |
| void spin_unlock(spinlock_t *lock) | Releases given lock |
| void spin_unlock_irq(spinlock_t *lock) | Releases given lock and enables local interrupts |
| void spin_unlock_irqrestore(spinlock_t *lock, unsigned long flags) | Releases given lock and restores local interrupts to given previous state |
| void spin_unlock_bh(spinlock_t *lock) | Releases given lock and enables bottom halves |
| void spin_lock_init(spinlock_t *lock) | Initializes given spinlock |
| int spin_trylock(spinlock_t *lock) | Tries to acquire specified lock; returns nonzero if lock is currently held and zero otherwise |
| int spin_is_locked(spinlock_t *lock) | Returns nonzero if lock is currently held and zero otherwise |





Semaphores

- User level:
 - Linux provides a semaphore interface corresponding to that in UNIX SVR4
- Internally:
 - implemented as functions within the kernel and are more efficient than user-visible semaphores
- Three types of kernel semaphores:
 - binary semaphores
 - counting semaphores
 - reader-writer semaphores
 - Multiple concurrent readers
 - Only one writer



Same as in Chapter 5



Linux Semaphores

| Traditional Semaphores | | | | |
|---|---|--|--|--|
| <pre>void sema_init(struct semaphore *sem, int count)</pre> | Initializes the dynamically created semaphore to the given count | | | |
| <pre>void init_MUTEX(struct semaphore *sem)</pre> | Initializes the dynamically created semaphore with a count of 1 (initially unlocked) | | | |
| <pre>void init_MUTEX_LOCKED(struct semaphore *sem)</pre> | Initializes the dynamically created semaphore with a count of 0 (initially locked) | | | |
| void down(struct semaphore *sem) | Attempts to acquire the given semaphore, entering uninterruptible sleep if semaphore is unavailable | | | |
| int down_interruptible(struct semaphore *sem) | Attempts to acquire the given semaphore, entering interruptible sleep if semaphore is unavailable; returns -EINTR value if a signal other than the result of an up operation is received. | | | |
| <pre>int down_trylock(struct semaphore *sem)</pre> | Attempts to acquire the given semaphore, and returns a nonzero value if semaphore is unavailable | | | |
| void up(struct semaphore *sem) | Releases the given semaphore | | | |
| Reader-Writer Semaphores | | | | |
| <pre>void init_rwsem(struct rw_semaphore, *rwsem)</pre> | Initalizes the dynamically created semaphore with a count of 1 | | | |
| <pre>void down_read(struct rw_semaphore, *rwsem)</pre> | Down operation for readers | | | |
| <pre>void up_read(struct rw_semaphore, *rwsem)</pre> | Up operation for readers | | | |
| <pre>void down_write(struct rw_semaphore, *rwsem)</pre> | Down operation for writers | | | |
| <pre>void up_write(struct rw_semaphore, *rwsem)</pre> | Up operation for writers | | | |





Summary

Deadlock:

- the blocking of a set of processes that either compete for system resources or communicate with each other
- blockage is permanent unless OS takes action
- may involve reusable or consumable resources
 - Consumable = destroyed when acquired by a process
 - Reusable = not depleted/destroyed by use

Dealing with deadlock:

- prevention guarantees that deadlock will not occur
- detection OS checks for deadlock and takes action
- avoidance analyses each new resource request



